

Taxicab Trajectory Classification using Deep Learning

Ashley M. Schuliger

Department of Computer Science, Worcester Polytechnic Institute, amschuliger@wpi.edu

Additional Keywords and Phrases: Taxi trajectories, Deep learning, Recurrent neural networks

1. INTRODUCTION

Taxis have been a major source of public transportations in large cities, such as New York City, for decades. However, the recent rise of rideshare companies, such as Uber and Lyft, has steadily decreased the popularity of taxis. In 2014, taxicabs serviced approximately 170,900,000 passengers with Uber only servicing 4,500,000 passengers. The scene changed drastically over three years as the approximate number of passengers served by Uber climbed to 159,900,000 and the number of passengers served by taxi-cabs declined to 125,500,000 [1]. This decrease in taxi usage is best explained by the wait time difference between each service. In a study by Forbes, passengers stated that they chose a ride-sharing service instead of a taxi most commonly due to the short wait times and travel times of the ride-share [2]. Due to this, it is imperative that the taxi industry find ways to decrease their overall wait time and travel time in order to regain the confidence of their customers. One way they could improve these statistics is by using their final destination to make decisions regarding their subsequent passengers. As a result of recent technology innovations, taxi cabs are equipped with GPS technology that allows the taxi dispatchers to better coordinate the passengers that taxis would serve by knowing the location of each taxi. However, GPS technology has no way of determining the final destination of a taxi driver, rendering the current GPS coordinates of a taxicab useless. Ride-share services, such as Uber and Lyft, have been able to clear this hurdle by allowing drivers to choose who their next passenger will be. This allows drivers to use their own final destination to determine which passenger they can service with the shortest wait time. While this framework makes it possible for Uber and Lyft to service passengers quickly, this framework is not possible for the taxi industry. While taxi drivers are employees of a company, Uber and Lyft drivers are independent contractors, which allows them to build their own schedule and choose their own passengers. Taxi drivers are given passengers from a taxi dispatcher and thus, cannot use their final destination to choose their next passenger. However, past

trajectories of these taxi cab drivers contain an overwhelming amount of information regarding their driving behaviors, including their speed along a route, their most frequently visited regions, and the average distance that they drive to serve a passenger. Due to this, it may be possible to use these driver behaviors to determine whether or not a trajectory should be assigned to a specific driver or not. In the subsequent sections, we propose the use of a deep neural network in order to predict the taxi driver based on a given trajectory. Our neural network is structured with two recurrent neural networks (RNN) that process seeking and serving sub-trajectories, and a linear activation function that takes in the output of these RNNs and five general features extracted from the complete trajectory. We provide the methodology used for building this neural network as well as the training and validation process that we used in order to obtain our model. We then train our model using cross validation to measure the accuracy. We also experiment with adjusting our hyperparameters in order to increase our accuracy further. Once trained, we test our complete neural network on a held out dataset and achieve an accuracy of 70%.

2. METHODOLOGY

In this section, we provide a methodology by which to classify and accurately predict the taxi driver based on their trajectory. We also describe the data processing and feature generation necessary for building a robust neural network for prediction.

A. Data Processing

Our dataset consists of a series of records where each record represents the GPS location of a taxi driver at the specific time in their trajectory. Each of these points is labeled with a 0 or 1 depending on whether the taxi cab is seeking or serving a passenger respectively. As such, this dataset contains long sequences of data points with the same serving or seeking label, indicating a subsequence of GPS coordinates along a seeking or serving sub-trajectory for the taxi driver. These sub-trajectories are important as they provide insight as to how a taxi driver behaves when seeking versus serving. Due to this, we chose to process the data into two separate lists: seeking and serving sub-trajectories. Each

element in these lists contain a list of GPS coordinates that represent a route that a taxi driver took to serve or seek out a passenger. This processing method allows us to make more sense of our data and take a closer look at the behavior of taxi drivers on route. Another important processing technique that was used included shifting our GPS coordinates into a grid layout. Longitude and latitude coordinates are precise to the degree by which no two data records are likely to share the same coordinates. Due to this, it was important to simplify our trajectory space to a 500 by 500 grid to reduce the precision of this data and provide a more general location for each data point. The grid size was set to 500 as this size generalizes the location of a point without losing a large degree of precision from the GPS coordinates. In order to bound this grid, we determined the maximum and minimum latitudes and longitudes for the entire dataset. Each data point was then assigned a cell with an x and y coordinate based on their latitude and longitude, where $0 \leq x \leq 500$ and $0 \leq y \leq 500$. These preprocessing techniques provided us with an insightful representation of our data. However, data cleaning was needed before preprocessing due to outliers in our GPS coordinates as well as abnormal lengths of subsequences found in our dataset. When calculating the bounds of our grid space, abnormal longitude and latitude coordinates skewed the size of our grid. Due to this, our grid space consisted of many sparse spaces and a small number of extremely dense spaces. As a result, individual cells in dense regions provided useless data as many points resulted in the exact same cell space, indicating that a taxi driver was not moving. Due to this, we chose to remove outliers in the data and limit the grid space to only the top 95% of GPS locations. This provided us with a grid space that better separated the majority of our data into cells that provided useful data. Along with this, our data contained short sequences of seeking and serving sub-trajectories where the length was less than 2. This represents a driver that only sought another passenger for a few seconds, which does not logically make sense. We chose to label these sequences as outliers and removed them from our dataset. Due to this, we concatenated the seeking or serving sub-trajectories on either side of this outlier to continue the sequence smoothly. Once we had pre-processed and cleaned our data, we chose to mold our data into the shape and size that we needed for our neural networks. In a subsequent section, we discuss our use of an RNN on our sub-trajectory data. In order to properly use such a neural network, our data needs to contain lists of consistent length. In our

case, we needed to input a list of sub-trajectories where each trajectory has the same length and each driver has the same number of sub-trajectories. We reconciled this by setting a constant for the number of sub-trajectories needed for each driver and another constant for the number of points needed in each sub-trajectory. We chose our constants based on their average values across all of the trajectory data we obtained. In terms of preprocessing, we deleted extra data points from the sub-trajectories and deleted sub-trajectories from a driver if the length was over the set limit. If the length of a sub-trajectory was below the constant, we filled the remaining length by duplicating the last data point in the original sub-trajectory. If the number of sub-trajectories for a driver was below the limit, we artificially added sub-trajectories of the limit length where each data point was located at $[-1, -1]$ on the grid. From here, we were able to use our dataset to effectively extract important features and build an accurate model.

B. Feature Generation

For feature generation, we focused on extracting features in two different ways: from the individual points and over the general trajectory. We chose to extract features for the individual points in order to gain a better understanding of how each taxi driver behaves at each time-step along a sub-trajectory. Specifically, we calculated the speed of the driver at each point as well as the hour within which the point was documented. Both of these features are important as they provide some insight as to how a driver moves during a sub-trajectory as different parts of the day. This information also allows us to take into account the type of traffic that a taxi driver may experience at certain times in a day, which can help us to better classify taxi drivers based on what type of route they take (i.e. one with more traffic, one with less traffic but a longer distance). These features are stored in the lists of seeking and serving sub-trajectories along with the cell coordinates of the data point. As for the scope of the entire trajectory, we wanted to extract features that described the general behavior of a taxi driver throughout the day. These features include the coordinates of the most visited cell, the average distance of a seeking sub-trajectory (before pre-processing), the average distance of a serving sub-trajectory (before pre-processing), and the number of passengers served in a trajectory. The most frequently visited cell provides us insight as to where the taxi driver is most likely to drive through or most likely to stay between passengers. The average distances show us the behavior of a driver in terms of whether they favor driving a shorter distance through

traffic or avoiding traffic by driving the extra miles and how far they are willing to drive to seek out a passenger. The number of passengers that a taxi driver serves in a trajectory is also important in that it shows us how efficient a taxi driver is compared to another. We use these extracted features in subsequent sections to build and train our deep neural network.

C. Network Structure

As mentioned in subsequent sections, our preprocessed data consists of two lists of subjectories (both seeking and serving) with various features included as well as a separate list of features for the general trajectory. Due to the nature of GPS coordinates, both lists of sub-trajectories are set-up in such a way that information can be extracted based on the sequence of points. Thus, we chose to use two recurrent neural networks to obtain and process more information from both of these sub-trajectories. From here, we wanted to process the output of both RNNs alongside the general trajectory features. In order to do this, we concatenated the RNN outputs to the feature set, and we fed those inputs into a simple linear activation function. We then fed these outputs into a softmax activation function in order to assign each output a driver classification. We chose this structure so that we would be able to extract as much information from the sequences of the trajectories as possible while also using the general features to classify these drivers. We also built a simple neural network that contains one linear activation function to classify each taxi driver, but as we will see in subsequent sections, our deep neural network far outperformed this simple neural network.

D. Training and Validation Process

When training our neural network, we used a combination of loss measurements and gradient descent techniques to obtain the optimal model. We use Cross Entropy Loss to determine the performance of the given model, and then we perform backwards propagation to update our weights accordingly. From here, we step forward once again using our gradient descent mechanism with the given learning rate. For gradient descent, we chose to use Pytorch's Adam gradient descent function. This version of gradient descent obtains an optimal model by performing the normal RMSProp gradient descent with momentum taken into account. This form of gradient descent performs well when there are multiple local minimum losses, as it allows for more exploration to find the global minimum. We regulate the number of times that gradient descent is

performed and the weights of the neural network are updated using a for loop with a number of epochs. In subsequent sections, we discuss how we chose the number of epochs and the learning rate for our gradient descent function.

3. EVALUATION & RESULTS

This section will cover the performance of our above mentioned neural network structure. We will first discuss the training and validation results using 5-fold cross validation of the original data set. Next, we evaluate the performance of our model against a baseline simple neural network. We also provide an explanation for how we chose the correct parameters for our resulting neural network.

A. Training & Validation Results

We start by providing the results of training our model in order to justify our chosen network structure. Our final model was trained using two RNNs with 4 input features and 2 outputs and a simple linear activation function with the RNN outputs concatenated to 5 general features. These outputs were fed into a softmax activation function in order to obtain the final results. In regards to training, we trained our model with 7 epochs and a learning rate of 0.01. Initially, we split our data into a training set and a testing set, but we wanted to try out different folds of our dataset. Thus, we performed 5-fold cross validation to provide us with a wider range of models to choose from. Figure 1 below shows the resulting training and testing accuracies after performing 5-fold cross validation.



Figure 1: Training and testing accuracies obtained through training our deep neural network

As shown in the figure above, our highest accuracy for both training and testing was achieved by our second validation fold with a training accuracy of 51.91% and a testing accuracy of 55.88%. In order to fully validate the performance of this model, we held out eight trajectories to perform a final validation. Using this optimal model, we achieved a 70% accuracy on this held out data, which implies a strong performance of our model.

B. Performance Comparison to Baseline

Based on the results above, this model provides a strong classification for taxi drivers based on their trajectory. In order to provide further proof of this, we compared our final results to the results produced by a baseline model. In our case, our baseline model is a simple neural network with only the five general features and none of the specific sequence data and features taken into account. As such, our simple neural network takes in 5 features and the outputs are then fed into a softmax activation function for the final classification. We performed 5-fold cross-validation on this baseline model to obtain classification results similar to those obtained in the previous section. Figure 2 below shows the results from using this baseline model.




Figure 2: Training and testing accuracies for the simple neural network without the use of RNNs

The figure above shows a poor performance in our baseline model, as the highest testing accuracy obtained with this model is 25.29%. The poor performance of this neural network provides justification for a deep neural network that takes into account information about the sequences of GPS coordinates along the driver's route. Thus, our model provides a more accurate prediction for the driver based on a trajectory than a simple neural network.

C. Hyperparameters

Once we had the structure for our neural network, we worked towards increasing our accuracy by adjusting the hyperparameters used in our training process. The hyperparameters that we chose to tune include the number of output layers in our activation functions, the learning rate for gradient descent, and the number of epochs used. For our RNN activation function, we were able to choose the number of outputs. Initially, we started with the same number of outputs as there were inputs. We then experimented by increasing and decreasing this value in order to analyze the effect it had on our accuracy. In doing this, we noticed that our optimal number of outputs was 2, which was lower than the number of inputs provided to the function. Along with this, we attempted to increase our accuracy by adding another simple linear activation function to our neural network. This linear activation function took in the five general features of the trajectories and the output was fed into the final linear activation function. Unfortunately, each iteration with this activation function decreased our accuracy, so we chose to not include it in our final neural network. For the learning rate, we started with a value of 0.05 and experimented with values above and below. As we increased this value, we noticed a sharp decrease in our accuracy, as our gradient descent was unable to train accurately with a larger learning rate. When we decreased our learning rate to 0.01, we found our optimal model. Thus, our final model uses a learning rate of 0.01 for gradient descent. For our epochs, we tested out different values of epochs to determine the optimal value for our model. In theory, as the number of epochs increases, the training error decreases; however, the testing error will decrease and then steadily increase. Due to this, it is important for us to test a wide range of epochs in order to find the optimal value where both our training and testing errors are low and our accuracies are high. In order to do this, we experimented with epoch values ranging from 1 to 15 and recorded the accuracies achieved with each one. Figure 3 below shows the resulting accuracies for each number of epochs.



Figure 3: Training and testing accuracies for various epoch values using our deep neural network

Unfortunately, the figure above does not follow the theory mentioned above. This could be due to a variety of reasons, and it may be useful to experiment with epoch values on different neural network models derived in previous sections. However, based on these results in the figure above, we will choose an epoch value of 7 since it provides the highest testing accuracy before it begins to decrease again. We are hesitant to choose 14, even though it achieves a higher accuracy, since a higher epoch value runs the risk of overfitting our model to the training dataset. As a result of these tests, we chose to adjust our model to have 2 outputs for our RNNs, a learning rate of 0.01 for our gradient descent, and 7 epochs for training.

4. CONCLUSION

In this paper, we proposed a deep neural network as a way to predict the taxi driver based on a trajectory. Our structure consisted of two recurrent neural networks with the outputs fed into one linear neural network and softmax. Training our model using Cross Entropy loss and Adam gradient descent aided in finding the optimal weights for our given neural network. Implementation on our training and testing set also validated the accuracy of our resulting model. We also provide a comparison of our model to a baseline model in order to prove the strong performance of our model. The proposed neural network provides strong predictions for the taxi driver given a trajectory. In future work, we plan to explore different features that we can generate from our original dataset in order to improve the performance of our neural network.

REFERENCES

1. (n.d.). Uber Revenue and Usage Statistics - BuildFire. Retrieved March 18, 2021, from <https://buildfire.com/uber-statistics/>
2. (2014, September 8). Uber, Lyft Cars Arrive Much Faster Than Taxis, Study Says - Forbes. Retrieved March 18, 2021, from <https://www.forbes.com/sites/ellenhuet/2014/09/08/uber-lyft-cars-arrive-faster-than-taxis/>